Module 6: Introduction to Metropolis

Rebecca C. Steorts

```r
knitr::opts_chunk$set(cache=TRUE)
```

Agenda

- Motivation
- Markov chain Monte Carlo (MCMC)
- Hard Discs in a Box Example
- Metropolis Algorithm
- Example Applied to Normal-Normal
- Practice Excerise (Hoff 10.3)

Intro to Markov chain Monte Carlo (MCMC)

Goal: sample from $f(x)$, or approximate $E_f[h(X)]$.

Recall that $f(x)$ is very complicated and hard to sample from.

How to deal with this?

1. What's a simple way?
2. What is another way?
3. What happens in high dimensions?

High dimensional spaces

- In low dimensions, rejection sampling works pretty well.
- But in high dimensions we run into problems.
- Why? It's hard to capture high dimensional spaces!

Figure 1: A high dimensional space (many images).

We turn to Markov chain Monte Carlo (MCMC).

Intuition

starting point

Figure 2: Example of a Markov chain and red starting point

Intuition

Figure 3: Example of a Markov chain and moving from the starting point to a high probability region.

What is Markov Chain Monte Carlo

- ▶ Markov Chain – where we go next only depends on our last state (the Markov property).
- ▶ Monte Carlo – just simulating data.

The Markov property

Suppose that we have just visited states $x_1, \ldots, x_{n-1}$. The Markov property says the following:

$$P(X_n = x_n \mid X_1 = x_1, \ldots, X_{n-1} = x_{n-1}) = P(X_n = x_n \mid X_{n-1} = x_{n-1}).$$

Why MCMC?

(a) the region of high probability tends to be "connected"
- ▶ That is, we can get from one point to another without going through a low-probability region, and

(b) we tend to be interested in the expectations of functions that are relatively smooth and have lots of "symmetries"
- ▶ That is, one only needs to evaluate them at a small number of representative points in order to get the general picture.

Advantages/Disadvantages of MCMC:

Advantages:

- ▶ applicable even when we can't directly draw samples
- ▶ works for complicated distributions in high-dimensional spaces, even when we don't know where the regions of high probability are
- ▶ relatively easy to implement
- ▶ fairly reliable

Disadvantages:

- ▶ slower than simple Monte Carlo or importance sampling (i.e., requires more samples for the same level of accuracy)
- ▶ can be very difficult to assess accuracy and evaluate convergence, even empirically

Hard Discs in a Box Example

Figure 4: Example of a phase diagram in chemistry.

Many materials have phase diagrams that look like the picture above.

Hard Discs in a Box Example

To understand this phenoma, a theoretical model was proposed:
Metropolis, Rosenbluth, Rosenbluth, and Teller, 1953

Figure 5: Example of N molecules (hard discs) bouncing around in a box.

Called hard discs because the molecules cannot overlap.

Hard Discs in a Box Example

Have an $X = (u, v)$ coordinate for each molecule.[1]

The total dimension of the space is $\mathbb{R}^{2N}$.

Figure 6: Example of N molecules (hard discs) bouncing around in a box.

---

[1] For the rest of the lecture, X will denote the two coordinate vectors of the molecule.

Hard Discs in a Box Example

$X \sim f(x)$ (Boltzman distribution).

Goal: compute $E_f[h(X)]$.

Since X is high dimensional, they proposed "clever moves" of the molecules.

Hard Discs in a Box Example

Metropolis algorithm: For iterations $i = 1, \ldots, n$, do:

1. Consider a molecule and a box around the molecule.
2. Uniformly draw a point in the box.
3. According to a "rule", you accept or reject the point.
4. If it's accepted, you move the molecule.

Uniformly = pick a point at random with equal probability to all other points in the box

Example of one iteration of algorithm

Consider a molecule and a box around the molecule.


Figure 7: This illustrates step 1 of the algorithm.

Example of one iteration of algorithm

Uniformly draw a point in the box.

Figure 8: This illustrates step 2 of the algorithm.

Example of one iteration of algorithm

According to a "rule", you accept or reject the point.

Here, it was accepted, so we move the molecule.

Figure 9: This illustrates step 3 and 4 of the algorithm.

Example of one iteration of algorithm

Here, we show one entire iteration of the algorithm.

Figure 10: This illustrates one iteration of the algorithm.

After running many iterations n (not N), we have an approximation for $E_f[h(X)]$, which is $\frac{1}{n}\sum_i h(X_i)$.

We will talk about the details later of why this is a "good approximation."

Some food for thought

We just covered the Metropolis algorithm (1953 paper).
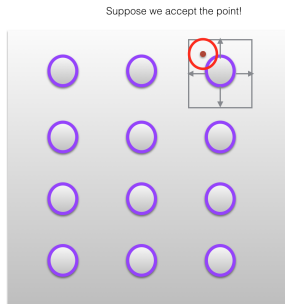
- ▶ We did not cover the exact procedure for accepting or rejecting (to come).
- ▶ Are the X_i's independent?
- ▶ Our approximation holds by The Ergodic Theorem for those that want to learn more about it.
- ▶ The ergodic theorem says: "if we start at a point x_o and we keeping moving around in our high dimensional space, then we are guaranteed to eventually reach all points in that space with probability 1."

Metropolis Algorithm

Setup: Assume pmf π on $\mathcal{X}$ (countable).

Have $f : \mathcal{X} \to \mathbb{R}$.

Goal:

a) sample/approximate from π

b) approximate $E_\pi[f(X)], X \sim \pi$.

The assumption is that π and or $f(X)$ are complicated!

Why things work!

Big idea and why it works: we apply the ergodic theorem.

"If we take samples $X = (X_0, X_1, \ldots,)$ then by the ergodic theorem, they will eventually reach π, which is known as the stationary distribution (the true pmf)."

Metropolis Algorithm

The approach is to apply the ergodic theorem.

1. If we run the Markov chain long enough, then the last state is approximately from π.

2. Under some regularity conditions,

$$\frac{1}{n} \sum_{i=1}^{n} f(X_i) \xrightarrow{a.s} E_\pi[f(X)].$$

Terminology

1. Proposal matrix = stochastic matrix.

   Let
   $$Q = (Q_{ab} : a, b \in \mathcal{X}).$$

   Note: I will use $Q_{ab} = Q(a, b)$ at times.

2. Note:
   $$\pi(x) = \tilde{\pi}(x)/z, z > 0.$$

   What is known and unknown above? (Think back to rejection sampling)

Metropolis Algorithm

Assume a symmetric proposal matrix Q. So, $Q_{ab} = Q_{ba}$.

Initialize a starting point $x_o \in X$.

Goal: run n total iterations of the algorithm to produce $(x_o, x_1, \ldots, x_n)$ samples.

This is useful as

$$\frac{1}{n} \sum_{i=1}^{n} f(x_i) \approx E_\pi[f(X)].$$

Metropolis Algorithm

For $i \in 1, 2, \ldots, n$:

- ▶ Sample a new data point x^* from $Q(x_i, x^*)$ if x is discrete, otherwise, $p(x^* \mid x_i)$.
- ▶ Sample u from Uniform$(0, 1)$.
- ▶ Let
$$r = \frac{\tilde{\pi}(x^*)}{\tilde{\pi}(x_i)}.$$
  If $u < r$, accept and $x_{i+1} = x^*$. accept the new data point
- ▶ Otherwise, reject and $x_{i+1} = x_i$. keep the current data point

Output: $x_o, x_1, \ldots, x_n$

Comment: r is the rule for uniformly drawing a point in the box (slide 18).

Metropolis Algorithm (Intuition)

How do we decide to accept or reject x^*?

In short, what is the intuition regarding how one would construct the Metropolis algorithm? [2]

Metropolis Algorithm (Intuition)

Suppose that we are considering to add a new data point x^* that is close to our current state x_i.

1. If $\tilde{\pi}(x^*) > \tilde{\pi}(x_i)$, then we would want to add a data point that is "close to" an existing data point in our set x_i.

2. However, if $\tilde{\pi}(x^*) < \tilde{\pi}(x_i)$, this case is not as clear cut. In short, we will want to accept/reject x^* in a logical way. (We will walk through this in detail below).

Metropolis Algorithm (Intuition)

Define
$$r = \frac{\tilde{\pi}(x^*)}{\tilde{\pi}(x_i)}.$$

1. Suppose $r > 1$.

Idea: Since x_i is in our data set already, we should include x^* since it has higher probability than x_i.

Procedure: Accept x^*, and update $x_{i+1} = x^*$.

Summary: We always accept x^* when $r > 1$. (This is the trivial case).

Metropolis Algorithm (Intuition)

2. Suppose $r < 1$.

Idea: The relative frequency values in our set equal to x^* compared to those equal to x_i will be r. This implies that for every instance of x_i we should only have a fraction of an instance of a x^* value.

Procedure: Set x_{i+1} to x^* or x_i with probability r and $1 - r$, respectively.

Summary: We will accept/reject the new data point x_{i+1} with probability r and $1 - r$.

Metropolis Algorithm (Intuition)

Reminder: We always accept in case 1, so we only need to consider case 2.

$$x_{i+1} = \begin{cases} x^* & \text{with probability} \quad \min(r, 1) \\ x_i & \text{with probability} \quad 1 - \min(r, 1) \end{cases}$$

We accomplish the sampling by drawing $u \sim \text{Uniform}(0, 1)$. Then $x_{i+1} = x^*$ if $u < r$. Otherwise, $x_{i+1} = x_i$.

Metropolis Algorithm (Correctness)

The Metropolis Algorithm is "correct" as one can show that it satisfies the conditions of the Ergodic Theorem. (Details are sketched along with definitions at the end of the slides for those that might be interested.)

Metropolis within a Bayesian setting

Goal: We want to sample from

$$p(\theta \mid y) = \frac{f(y \mid \theta)\pi(\theta)}{m(y)}.$$

Typically, we don't know $m(y)$.

The notation is a bit more complicated, but the set up is the same.

We'll approach it a bit differently, but the idea is exactly the same.

Building a Metropolis sampler

We know $\pi(\theta)$ and $f(y \mid \theta)$, so we can can draw samples from these.

Our notation here will be that we assume parameter values $\theta_1, \theta_2, \ldots, \theta_s$ which are drawn from $\pi(\theta)$.

We assume a new parameter value comes in that is θ^*.

Building a Metropolis sampler

Similar to before we assume a symmetric proposal distribution, which we call $J(\theta^* \mid \theta^{(s)})$.

▶ What does symmetry mean here? $J(\theta_a \mid \theta_b) = J(\theta_b \mid \theta_a)$.

▶ That is, the probability of proposing $\theta^* = \theta_a$ given that $\theta^{(s)} = \theta_b$ is equal to the probability of proposing $\theta^* = \theta_b$ given that $\theta^{(s)} = \theta_a$.

▶ Symmetric proposals include:

$$J(\theta^* \mid \theta^{(s)}) = \text{Uniform}(\theta^{(s)} - \delta, \theta^{(s)} + \delta)$$

and

$$J(\theta^* \mid \theta^{(s)}) = \text{Normal}(\theta^{(s)}, \delta^2).$$

Metropolis algorithm

The Metropolis algorithm proceeds as follows:

1. Sample $\theta^* \sim J(\theta \mid \theta^{(s)})$.

2. Compute the acceptance rule (r):

$$r = \frac{p(\theta^*|y)}{p(\theta^{(s)}|y)} = \frac{p(y \mid \theta^*)p(\theta^*)}{p(y \mid \theta^{(s)})p(\theta^{(s)})}.$$

3. Let

$$\theta^{(s+1)} = \begin{cases} \theta^* & \text{with prob min(r,1)} \\ \theta^{(s)} & \text{otherwise.} \end{cases}$$

Remark: Step 3 can be accomplished by sampling $u \sim \text{Uniform}(0,1)$ and setting $\theta^{(s+1)} = \theta^*$ if $u < r$ and setting $\theta^{(s+1)} = \theta^{(s)}$ otherwise.

Exercise: Convince yourselves that step 3 is the same as the remark!

A Toy Example of Metropolis

Let's test out the Metropolis algorithm for the conjugate Normal-Normal model with a known variance situation.

$$X_1, \ldots, X_n \mid \theta \overset{iid}{\sim} \text{Normal}(\theta, \sigma^2)$$
$$\theta \sim \text{Normal}(\mu, \tau^2).$$

Recall that the posterior of $\theta \mid X_1, \ldots, X_n$ is $\text{Normal}(\mu_n, \tau_n^2)$, where

$$\mu_n = \bar{x} \frac{n/\sigma^2}{n/\sigma^2 + 1/\tau^2} + \mu \frac{1/\tau^2}{n/\sigma^2 + 1/\tau^2}$$

and

$$\tau_n^2 = \frac{1}{n/\sigma^2 + 1/\tau^2}.$$

Toy Example

In this example: $\sigma^2 = 1$, $\tau^2 = 10$, $\mu = 5$, $n = 5$, and

$$x = (9.37, 10.18, 9.16, 11.60, 10.33).$$

For these data, $\mu_n = 10.03$ and $\tau_n^2 = 0.20$.

Note: this is a toy example for illustration.

Toy example

We need to compute the acceptance rule r.

$$r = \frac{p(\theta^*|x)}{p(\theta^{(s)}|x)} \tag{1}$$

$$= \frac{p(x|\theta^*)p(\theta^*)}{p(x|\theta^{(s)})p(\theta^{(s)})} \tag{2}$$

$$= \left(\frac{\prod_i \mathsf{dnorm}(x_i, \theta^*, \sigma)}{\prod_i \mathsf{dnorm}(x_i, \theta^{(s)}, \sigma)} \right) \left(\frac{\mathsf{dnorm}(\theta^*, \mu, \tau)}{\mathsf{dnorm}(\theta^{(s)}, \mu, \tau)} \right) \tag{3}$$

Toy example

In many cases, computing the rule r directly can be numerically unstable, however, this can be modified by taking $\log r$.

This results in

$$
\begin{aligned}
\log r = \sum_i & \left[\log \mathrm{dnorm}(x_i, \theta^*, \sigma) - \log \mathrm{dnorm}(x_i, \theta^{(s)}, \sigma) \right] \\
& + \left[\log \mathrm{dnorm}(\theta^*, \mu, \tau) \right] - \log \mathrm{dnorm}(\theta^{(s)}, \mu, \tau).
\end{aligned}
$$

Then a proposal is accepted if $\log u < \log r$, where u is sampled from the Uniform(0,1).

Toy example

We generate 500 iterations of the Metropolis algorithm starting at
$\theta^{(0)} = 0$ and using a normal proposal distribution, where

$$\theta^{(s+1)} \sim \text{Normal}(\theta^{(s)}, 1).$$

We will then generate 10,000 iterations since 500 will not be
sufficient.

Figure~12 shows a traceplot for this run as well as a histogram for
the Metropolis algorithm compared with a draw from the true
normal density.

Traceplot

What is a traceplot? A traceplot is a convergence diagnostic.

It's a plot of the parameter of interest versus the number of MCMC iterations.

What does it tell us?

1. It can tell us when we have not run our MCMC long enough.
2. It can tell us when we migth be in a situation of multimodality (we will see this in future lectures).
3. If the plots looks stable, then it tells us that we do not see anything to warrant issues with a lack of convergence of the chain.

What does it not tell us?

It cannot tell us that our MCMC has converged!

Code

```
# setting values
set.seed(1)
s2<-1
t2<-10
mu<-5;
n<-5

# defining the data
x<-c(9.37, 10.18, 9.16, 11.60, 10.33)
# mean of the normal posterior
mu.n<-( mean(x)*n/s2 + mu/t2 )/( n/s2+1/t2)
# variance of the normal posterior
t2.n<-1/(n/s2+1/t2)
```

Code (continued)

```
##S = total num of simulations
theta<-0 ; delta<-1 ; S<-1000 ; THETA<-NULL

for(s in 1:S){
## simulating our proposal
theta.star<-rnorm(1,theta,sqrt(delta))
##taking the log of the ratio r
log.r<-( sum(dnorm(x,theta.star,sqrt(s2),log=TRUE)) +
dnorm(theta.star,mu,sqrt(t2),log=TRUE) ) -
( sum(dnorm(x,theta,sqrt(s2),log=TRUE)) +
dnorm(theta,mu,sqrt(t2),log=TRUE) )
if(log(runif(1))<log.r) {
  theta<-theta.star
}

THETA<-c(THETA,theta) ##updating THETA

}
```

Code (continued)

```r
length(THETA)
```

```
## [1] 50
```

```r
head(THETA)
```

```
## [1] 0.000000 1.329799 2.925080 2.925080 3.412509 3.41250(
```

Code (continued)

```
## pdf
##   2
```

Code (continued)

```
## pdf
##   2
```

Traceplot and Histogram

Figure 11: Left: trace plot of the Metropolis sampler. Right: Histogram versus true normal density for 500 iterations.

Traceplot and Histogram

Figure 12: Left: trace plot of the Metropolis sampler. Right: Histogram versus true normal density for 10,000 iterations.

Traceplot

Given that we have looked at $n = 500$ iterations of the Metroplis sampler, does it seem that our approximation is a good one?

What about $n = 10,000$ iterations?

Questions you should be able to answer!

- ▶ What is the goal of Metropolis?
- ▶ What is known and unknown?
- ▶ What are good proposals?
- ▶ What does the ergodic theorem say in words?
- ▶ Are good proposals always easy to choose?
- ▶ When would we use Metropolis over Importance sampling and Rejection sampling?
- ▶ What is a simple diagnostic of a Markov chain?
- ▶ Are we guaranteed convergence of the Markov chain?

Detailed Takeways

- ▶ Monte Carlo is not sufficient for high dimensions
- ▶ Intuitive of MCMC
- ▶ What is Monte Carlo
- ▶ What is a Markov Chain
- ▶ What is the memoryless property
- ▶ Why do we use MCMC
- ▶ Advantages versus Disdavantages of MCMC
- ▶ Metropolis et al. (1953) original paper and understanding this case study
- ▶ Ergodic Theorem
- ▶ Metropolis Algorithm (from the 1953 paper)
- ▶ Metropolis Algorithm in a Bayesian setting
- ▶ Normal-Normal Example using Metropolis
- ▶ Traceplots: What they tells us and what they don't tells us
- ▶ How long should you run an MCMC chain
- ▶ How you would walk through a case study of metropolis but on the normal-uniform model (very similar to the in class study)

Information about Markov chains

If you're wondering what a Markov chain and a problem looks like in a discrete space, take a look at this video:
https://www.youtube.com/watch?v=i3AkTO9HLXo

If you're wanting to read more about Markov chains, see
http://faculty.washington.edu/yenchic/18A_stat516/Lec3_DTMC_p1.pdf

Ergodic Theorem

Suppose $(X_o, X_1, \ldots, X_n)$ is an irreducible (time-homoegenous) discrete Markov Chain with stationary distribution π.

Then

$$\frac{1}{n} \sum_{i=1}^{n} f(X_i) \overset{a.s.}{\to} E_f[X]$$

as $n \to \infty$ where $X \sim \pi$. Also, $f : \mathcal{X} \to R$ is a bounded function.

If the Markov Chain is aperiodic, then

$$P(X_n = x \mid X_o = x_o) \to \pi(x)$$

as $n \to \infty$ for all $x, x_o \in \mathcal{X}$.

Notation

Let $(X_i) = (X_o, X_1, \ldots, X_n)$.

Time-Homogeneous Markov Chain

A Markov Chain (X_i) is time homogeneous if

$$P(X_{i+1} = b \mid X_i = a) = T_{ab}$$

for all i, for all $a, b \in \mathcal{X}$, and for matrix T.

The matrix T is called the transition matrix (which is a stochastic matrix).

Stationary Distribution

A pmf π defined on $\mathcal{X}$ is a stationary distribution (with respect to the transition kernel) if

$$\pi T = \pi.$$

Another way to write this is as follows:

$$\sum_{a \in \mathcal{X}} \pi_a T_{ab} = \pi_b$$

for all $b \in \mathcal{X}$.

Irreducible

A Markov Chain (X_i) is irreducible if for all $a, b \in \mathcal{X}$ there exists $t \geq 0$ such that

$$P(X_t = b \mid X_o = a) > 0.$$

Aperiodic

An irreducible Markov chain (X_i) is aperiodic if for all $a \in \mathcal{X}$

$$gcd\{P(X_t = a \mid X_o \mid a) > 0\} = 1.$$

greatest common divisor $= gcd$

Metropolis (Correctness)

We need to prove that the ergodic theorem holds. (I'll provide a sketch of the proof, which requires knowledge of measure theory/real analysis).

The transition kernel can be written as follows:

$$T(x, x_i) = P(X_{i+1} = x \mid X_i = x_i) \tag{4}$$
$$= p(x \mid x_i) \times p(\text{accept } x \mid x, x_i) \tag{5}$$
$$= Q(x_i, x) min(1, \tilde{\pi}(x)/\tilde{\pi}(x_i)) \tag{6}$$

as long as $x_i \neq x$.

Proposition: Detailed Balance

X_i has stationary distribution π.

Proof: Need to verify that $\pi_a T_{ab} = \pi_b T_{ba}$.

a. If $a = b$, the above trivially holds.
b. If $a \neq b$,

$$\pi(a)T(a,b) = \pi(a)Q(a,b)\min(1, \tilde{\pi}(x)/\tilde{\pi}(x_i)) \qquad (7)$$

$$= Q(a,b)\min(\pi(a),\pi(b)) \quad \text{symmetric about a,b} \qquad (8)$$

$$= \pi(b)T(b,a) \qquad (9)$$

Thus, detailed balance is satisfied.

Result: Aperiodity and Irreduciblity

Assume $\pi(x) > 0$ for all $x \in \mathcal{X}$.

1. If Q is irreducible, then T is irreducible.
2. If Q is aperiodic, then T is aperiodic.

Proposition

If $\pi(x) > 0$ for all $x \in \mathcal{X}$ and Q is irreducible then

$$\frac{1}{n} \sum_{i=1}^{n} f(X_i) \overset{a.s.}{\to} E_f[X]$$

as $n \to \infty$.

Also, if Q is aperiodic, then

$$P(X_n = x \mid X_o = x_o) \to \pi(x)$$

as $n \to \infty$ for all $x \in \mathcal{X}$.